Melanie Subbiah

m.subbiah@brown.edu
melaniesubbiah.github.io

Academic Appointments

2026 – *present* **Brown University** Providence, RI
Postdoctoral Researcher in Computer Science
Supervisor: Ellie Pavlick

Education

2020 – 2026 **Ph.D. Computer Science**, Columbia University
Thesis: *Interpretation and Interpretability in Subjective Domains*
Advisor: Kathleen McKeown

2020 – 2022 **M.S. Computer Science**, Columbia University

2013 – 2017 **B.A. Computer Science**, Williams College
Thesis: *Using Text Abstraction and LSTM Language Models for Domain-Independent Narrative Generation*
Advisor: Andrea Danyluk

Publications

* indicates equal contribution

Conference and Journal Papers

[NeurIPS 25] Hall, Z., **Subbiah***, **M.**, Zollo*, T., McKeown, K., & Zemel, R. (2025). Guiding LLM Decision-Making with Fairness Reward Models. In *Proceedings*, Neural Information Processing Systems.

[INLG 25] Limpijankit, M., Chen, Y., **Subbiah**, **M.**, Deas, N., & McKeown, K. (2025). Counterfactual Simulatability of LLM Explanations for Generation Tasks. In *Proceedings*, International Natural Language Generation.

[EMNLP 25] **Subbiah**, **M.**, Mishra, A., Grace Kim, L. T., Durrett, G., & McKeown, K. (2025). Is the Top Still Spinning? Evaluating Subjectivity in Narrative Understanding. In *Proceedings*, Empirical Methods in Natural Language Processing.

[TACL 24] **Subbiah**, **M.**, Zhang, S., Chilton, L. B., & McKeown, K. (2024). Reading Subtext: Evaluating Large Language Models on Short Story Summarization with Writers. In *Transactions of the Association for Computational Linguistics*, MIT Press (Best Paper Award).

[EMNLP 24] **Subbiah***, **M.**, Ladhak*, F., Mishra, A., Adams, G., Chilton, L. B., & McKeown, K. (2024). STORYSUMM: Evaluating Faithfulness in Story Summarization. In *Proceedings*, Empirical Methods in Natural Language Processing.

[ACL 23] Storek, A., **Subbiah**, **M.**, & McKeown, K. (2023). Unsupervised Selective Rationalization with Noise Injection. In *Proceedings*, Association for Computational Linguistics.

[ACL 23] Wang, G., Chillrud, L., Harwood, K., Ananthram, A., **Subbiah**, **M.**, & McKeown, K. (2023). Check-COVID: Fact-Checking COVID-19 News Claims with Scientific Evidence. In *Findings*, Association for Computational Linguistics.

[EMNLP 22] Levy, S., Allaway, E., **Subbiah**, **M.**, Chilton, L., Patton, D., McKeown, K., & Wang, W. (2022). SafeText: A Benchmark for Exploring Physical Safety in Language Models. In *Proceedings*, Empirical Methods in Natural Language Processing.

[EMNLP 22] Mei, A., Kabir, A., Levy, S., **Subbiah**, **M.**, Allaway, E., Judge, J., … Wang, W. (2022). Mitigating Covertly Unsafe Text within Natural Language Systems. In *Findings*, Empirical Methods in Natural Language Processing.

[NeurIPS 20] Brown*, T., Mann*, B., Ryder*, N., **Subbiah***, **M.**, Kaplan, J. D., Dhariwal, P., … Amodei, D. (2020). Language Models are Few-Shot Learners. In *Proceedings*, Neural Information Processing Systems (Outstanding Paper Award).

[PLoSOne 17] Nygaard, U., Li, Z., Palys, T., Jackson, B., **Subbiah**, **M.**, Malipatlolla, M., … Nadeau, K. (2017). Cord blood T cell subpopulations and associations with maternal cadmium and arsenic exposures, PLoS One.

Refereed Workshops

[AI & Self NeurIPS 26] Wei, M., **Subbiah**, **M.**, Lee, A., Dahmani, A., Edwards, D., Edwards, H., & Pavlick, E. (2026). *Ownership in AI-Assisted Everyday Tasks*. AI and the Self Workshop at NeurIPS.

[NewSumm EMNLP 25] Gupta, M., Varimalla, N. R., Deas, N., **Subbiah**, **M.**, & McKeown, K. (2025). *AdvSumm: Adversarial Training for Bias Mitigation in Text Summarization*. New Frontiers in Summarization Workshop at EMNLP.

[ARP Symposium 25] Mayukha, A., **Subbiah**, **M.**, Guzman, A., Jitklongsub, S., & McAdams, D. (2025). *Methodological Approaches to the Analysis of Life Story Interviews*. Symposium at the Association for Research in Personality conference (ARP).

[WASSA ACL 23] **Subbiah***, **M.**, Bhattacharjee*, A., Hua, Y., Kumarage, T., Liu, H., & McKeown, K. (2023). *Detecting Harmful Agendas in News Articles*. Workshop on Computational Approaches to Subjectivity, Sentiment & Social Media Analysis at ACL.

[IWSS ICSWM 21] **Subbiah**, **M.**, & McKeown, K. (2021). *Understanding Identity Signalling in Persuasive Online Text*. International Workshop on Social Sensing at ICWSM.

[WiML NeurIPS 18] Maher, M., **Subbiah**, **M.**, & Apostoloff, N. (2018). *Cascaded Dataset QA*. Women in Machine Learning Workshop at NeurIPS.

[WiML NeurIPS 18] **Subbiah**, **M.**, lesser, J., & Apostoloff, N. (2018). *Augmenting Training Data with Simulated Images*. Women in Machine Learning Workshop at NeurIPS.

Preprints and Papers Under Review

[*in submission* 26] Gong, Z., Huang, Y., Deas, N., **Subbiah**, **M.**, & Hirschberg, J. (2026). *Right Destination, Wrong Journey: Synthetic Counseling Dialogues Miss Human Interactional Dynamics*.

[*in submission* 26] Gong, Z., Huang, Y., Liang, W., Deas, N., **Subbiah**, **M.**, McKeown, K., & Hirschberg, J. (2026). *Empath: Tracing Multi-Level Emotion Dynamics in Crisis Counseling Dialogues*. arXiv.

[arXiv 26] Mian*, H., **Subbiah***, **M.**, Marcus, S., Shaalan, N., & McKeown, K. (2026). *Computational Representations of Character Significance in Novels*. arXiv.

[arXiv 26] **Subbiah**, **M.**, Mian, H., Deas, N., Mayukha, A., McAdams, D. P., & McKeown, K. (2026). *Whose Story Gets Told? Positionality and Bias in LLM Summaries of Life Narratives*. arXiv.

[*in submission* 26] **Subbiah**, **M.**, Tsai, C.-E. J., & McKeown, K. (2026). *ToMStories: Interpreting Theory-of-Mind in Synthetic Narratives*. arXiv.

[*in submission* 26] **Subbiah***, **M.**, Hall*, Z., & McKeown, K. (2026). *On the Limits of Steering Vectors for Preference-Aligned Generation*. arXiv.

Patents

[patent 22] **Subbiah, M.**, Lesser, J., & Apostoloff, N. (2022). *Training with Simulated Images*. No. US 11,256,958 B1, Apple, Inc.

Industry Research Experience

2022 (summer) **Meta**, New York, NY
Mentor: Ron Weiss
Deep learning to interpret EMG data for human-computer neural interfaces.

2019 – 2020 **OpenAI**, San Francisco, CA
Mentors: Alec Radford, Dario Amodei, Daniela Amodei
Evaluation suite for GPT-3, co-first-author on the GPT-3 paper.

2017 – 2019 **Apple**, Cupertino, CA
Mentor: Nicholas Apostoloff
Reward function design for autonomous systems, domain randomization for sim-to-real transfer in computer vision, and effective QA methods for crowdsourced annotated data.

Grants and Fellowships

2026 Characterizing and Intervening on Behavioral Drift in Long Contexts
Goodfire, Brown University, $25,000

2023 Narrative Summarization
Amazon CAIT, Columbia University, PhD Fellowship

2020 NLP with Social Impact
Office of the President and SEAS, Columbia University, PhD Fellowship

Honors and Awards

2026 PhD Service Award, *Columbia University Department of Computer Science*

2024 Best Paper, *TACL*

Outstanding Reviewer, *EMNLP*

Most Influential Scholar Award Honorable Mention, *AMiner AI 2000*

2023 Top 100 AI Achievements from 1943-2021, *BenchCouncil AI100*

2021 Honorable Mention, *NSF Graduate Research Fellowship Program*

2020 Outstanding Paper, *NeurIPS*

2017 Highest Honors & Best Colloquium Presentation, *Williams College Computer Science*

Shorty Story Writing Honorable Mention, *Williams College Benjamin B. Wainwright Prize*

Teaching

Instructor

2023 Discrete Mathematics, *Columbia University*
Course Quality Overall: 4.2/5.0, Instructor Quality Overall: 4.5/5.0

Teaching (continued)

Teaching Assistant

2023 — Natural Language Generation and Summarization, *Columbia University*
Instructor: Kathleen McKeown
TA Quality Overall: 4.7/5.0

2022 — Natural Language Generation and Summarization, *Columbia University*
Instructor: Kathleen McKeown
TA Quality Overall: 4.9/5.0

2014 – 2017 — Introductory Computer Science and/or Data Structures, *Williams College*
Instructors: William Lenhart, Thomas Murtagh, Morgan McGuire

Guest Lectures

2025 — Computing in Context, *Columbia University*
Instructor: Debasmita Bhattacharya

2024 — Global Teaching Labs - Uruguay, *MIT & University of Montevideo*
Instructor: Yi-Tong Tseo

2023 — Computational Journalism, *Columbia University*
Instructor: Mark Hansen

Natural Language Generation and Summarization, *Columbia University*
Instructor: Kathleen McKeown

2022 — Natural Language Generation and Summarization, *Columbia University*
Instructor: Kathleen McKeown

Speaking

Invited Seminars

2024 — How Did We Get Here?: The Rise of LLMs and the Problem of Evaluation, *Williams College*

2020 — Language Models are Few-Shot Learners, *New York University*; *Stanford University*; *Columbia University*; *New York University*; *Philosophy and Machine Learning Conference*; *G-Research*

Language Models, *Williams College*

Conference Presentations

2026 — Optimism in AI: A Human Problem, *Hope and Optimism* (talk)

2025 — Is the Top Still Spinning? Evaluating Subjectivity in Narrative Understanding, *EMNLP* (poster)

2024 — StorySumm: Evaluating Faithfulness in Story Summarization, *EMNLP* (talk)

Reading Subtext: Evaluating LLMs on Short Story Summarization with Writers, *EMNLP* (poster)

2023 — Detecting Harmful Agendas in News Articles, *Computational Approaches to Subjectivity, Sentiment & Social Media Analysis Workshop @ ACL* (talk)

2021 — Understanding Identity Signalling in Persuasive Online Text, *International Workshop on Social Sensing @ ICWSM* (talk)

2018 — Augmenting Training Data with Simulated Images, *Women in Machine Learning Workshop @ NeurIPS* (poster)

Outreach and Media

2024 — Research Overview: Language Models, *Stuyvesant High School* (research class)

Speaking (continued)

2023	ChatGPT for Educators, *Academy for Teachers* (Master Class for 25 high school teachers and librarians)
	Artificial: Episode 2, Selling Out, *The Wall Street Journal* ("The Journal" podcast)
	NLP with ChatGPT, *Pearson Publishing* (AI Catalyst Conference)
	ChatGPT: What is it, how does it work, and what's next?, *Columbia University* (Council of Deans; faculty)
2022	GPT-3 for Natural Language Processing, *SuperDataScience* (podcast)
2021	5 Levels of Difficulty: Machine Learning, *Wired Magazine* (YouTube video)
2017	Phi Beta Kappa Student Speaker, *Williams College* (Commencement)

Mentoring

Brown University

Alexander Kaneko (Undergrad)

Columbia University

Haaris Mian (Masters) → Cambridge DAMTP PhD
Marvin Limpijankit (Masters) → INLG paper, Columbia CS PhD
Zara Hall (Undergrad/Masters) → NeurIPS paper, Columbia CS PhD
Mukur Gupta (Masters) → NewSumm workshop paper
Nikhil Varimalla (Masters) → NewSumm workshop paper
Jacklyn Tsai (Undergrad) → Senior Thesis
Akankshya Mishra (Masters) → two EMNLP papers, Google Gemini MLE
Sean Zhang (Masters) → TACL paper (Best Paper Award)
Kate Harwood (Masters) → ACL Findings paper
Adam Štorek (Undergrad/Masters) → ACL paper, Columbia CS PhD
Bobby Yilun Hua (Undergrad) → WASSA workshop paper, Cornell CS PhD

Service

Reviewing

2026	ACL, EMNLP, NAACL
2025	ACL, COLM, ARR, ICLR
2024	ACL, EMNLP, COLING, ARR
2022	EMNLP, In2Writing Workshop @ ACL
2019	Bay Area Machine Learning Symposium

Organizing

2026	Co-organizer, *AI and Mental Health Consortium*, ARIA Institute
2021	Co-organizer, *Workshop on Enormous Language Models*, ICLR

Department and Community

2023	Panelist, *Graduate Council info panel on PhD programs*, Columbia Engineering
2022	Reviewer, *Pre-application Review for underrepresented candidates*, Columbia University
	Mentor, *CS undergraduate buddy program*, Williams College

Service (continued)

2021 – 2025	Organizer, *Lab retreat*, Columbia University
2020	Reviewer, *Pre-application Review for underrepresented candidates*, Columbia University
	Speaker, *"Demystifying the Dissertation" talk series for undergrads*, Columbia University
2016 – 2017	Leader, *Underrepresented Identities in CS group*, Williams College

Consulting

2024	Advising on technical requirements for customers of a green data center, *BuildAI*
	Advising on models for educational children's story generation, *Nookly*
2023	Red teaming for GPT-4, *OpenAI*
	Technical reviewer for Quick Start Guide to Large Language Models, *Pearson Publishing*